

## Q1 Directions

6 Points

For this midterm you are allowed the use of your two-page cheat sheet, but no resources on the Internet. It is an 80-minute exam.

Clarification questions can be posted **privately only** to Piazza, and will be responded to there.

Before you submit your exam, please consult the Exam Clarifications here:

[https://docs.google.com/document/d/1lBhXtt5\\_0jd1g7eg7wKyGHRz20Blk-5AXMtuqSBryx0/edit?usp=sharing](https://docs.google.com/document/d/1lBhXtt5_0jd1g7eg7wKyGHRz20Blk-5AXMtuqSBryx0/edit?usp=sharing).

☒ I've read this and understand the rules.

## Q2 Capturing data

3 Points

Which can capture more information overall:

- ☐ relations
- ☐ graphs
- ☐ JSON / hierarchical data
- ☒ they are equivalent

## Q3 Query operations

4 Points

Which of the following are true (select all that apply):

- ☒ the results of an inner join are a subset of the results of an outer join
- ☐ the main operator for combining different rows from similar tables is the outer join
- ☐ the *applymap* operator is used to combine tables
- ☒ a filter or select operation can be performed in parallel across each tuple

## Q4 Knowledge representation

3 Points

If we are given a conceptual class of entities  $C$  and a subclass of  $C$  called  $S$ , we know that:

- ☒ every member of class  $S$  must also be a member of class  $C$
- ☐ class  $C$  represents a subset of class  $S$
- ☐ there is no formal relationship between  $S$  and  $C$
- ☐ class  $C$  has every property in class  $S$

## Q5 NoSQL

3 Points

Choose the *best* answer from the following.

- ☐ NoSQL databases cannot store nested objects.
- ☒ NoSQL databases typically have limited notions of consistency

☒ NoSQL databases typically have limited notions of consistency.

☐ SQL databases are non-relational.

☐ SQL databases are key-value based.

## Q6 Integration

3 Points

For the strings "danger" and "roger," how many 3-grams will be in common?

☐ 1

☐ 2

☒ 3

☐ 4

☐ 6

## Q7 Computer architecture

3 Points

Given a list of 768 single-byte characters, an iteration through the characters in the array, and a cache size of 256B, we can amortize memory accesses by up to a factor of:

☐ 3

☐ 16

☒ 256

☐ 768

## Q8 Sharding

4 Points

Give the **best** answer here. Given a table  $R(A, B, C)$  with 10M entries, if we use hash-based sharding on attribute  $A$  in this table, we can support up to how many worker nodes running in parallel:

- ☐ 1,000
- ☐ 1M
- ☐ 10M
- ☒ number of unique values of  $A$
- ☐ number of unique tuples in  $R$

## Q9 Queries

3 Points

Suppose we have a database of students, courses, and enrollments. If we want to know the number of *unique last names* of students who are enrolled in courses, we can:

- ☐ left outerjoin all students with enrollments, count the number of (distinct) student last names
- ☒ inner join all students with enrollments, count the number of (distinct) student last names
- ☐ union all students with enrollments, count the number of (distinct) student last names
- ☐ this query is not possible to express in SQL, Pandas, or relational operators

## Q10 Statistical tests

3 Points

Suppose we apply a statistical test, it indicates that the probability of the null hypothesis is under our  $\alpha$  level, and we decide that our candidate hypothesis holds. If, in reality our candidate hypothesis is wrong, then:

- ☒ we have a false positive
- ☐ we have a false negative
- ☐ we used the wrong test for the distribution
- ☐ we need to use the Bonferroni correction

## Q11 Indexing

4 Points

Indexing a relation or dataframe  $R(A, B, C)$  by attribute  $A$  helps with (check all that apply):

☒ joining on attribute  $A$

☒ filtering on attribute  $A$

☒ projecting attribute  $A$

☒ grouping by attribute  $A$

☐ projecting attributes  $(A, B)$

## Q12 Breadth-first search

3 Points

Which does less overall work, for most real-world graphs:

- ☒ sequential (centralized) breadth-first search
- ☐ parallel breadth-first search
- ☐ they do exactly the same amount of work
- ☐ breadth-first search cannot be parallelized

## Q13 Partitioning data

3 Points

Which of the following are likely to make it *more difficult* to balance our workload across machines, in a sharded query processing setting:

- ☐ dataframes are of different sizes from one another
- ☐ certain values occur frequently, within an attribute used in selection operations
- ☒ certain values occur frequently, within a join key
- ☐ all values within a grouping key appear exactly once

## Q14 Link analysis

3 Points

In the PageRank algorithm with the decay factor added, the total PageRank in the system should, across iterations:

- ☐ increase
- ☐ decrease
- ☒ remain constant

☐ any of the above, it depends on the structure of the graph

## Q15 Matrices for machine learning

3 Points

When we convert from a dataframe  $R$  with  $k$  attributes and  $n$  rows to a matrix  $M$  suitable for machine learning, how many columns can the matrix be? Consider cases where the dataframe is not purely numeric.

☐  $k$

☒  $k \cdot n$

☐  $n$

☐  $n^2$

## Q16 Distributed processing

6 Points

Name two relational operations in Spark (dataframes or SQL) require the system to shard on a particular key.

### Q16.1 Relational operator 1

3 Points

Your first answer:

GROUP BY

## Q16.2 Relational operator 2

3 Points

Your second answer:

JOIN

## Q17 Speedups

8 Points

Name two types of **data structures** we (Pandas, Spark, or the programmer) can take advantage of to speed up dataframe or query processing.

### Q17.1 Data structure 1

4 Points

Your first answer:

Dictionary

### Q17.2 Data structure 2

4 Points

Your second answer:

B+ Tree

## Q18 Long-form query

11 Points

Suppose we have a dataset about classes, and we are curious how



often we have multiple students **with the same first name** enrolled in a class. The tables are *Students*(*id*, *first*, *last*) and *Enrolled*(*studentid*, *courseid*).

The data is available in Spark dataframes. The *Students* table (dataframe) is sharded by *id* and the *Enrolled* table is not sharded by any particular key (i.e., tuples are randomly partitioned across our worker nodes).

### Q18.1 The query

6 Points

Write a query to find all such results (either name, class, count; or name, count), using Spark or Pandas operations or SQL, or pseudocode that gives enough detail to understand which operators are being performed.

```
ans = spark.sql("""SELECT Students.first, Enrolled.courseid,
count(*)
FROM Students JOIN Enrolled on Students.id =
Enrolled.studentid
GROUP BY Students.first, Enrolled.courseid""")
```

### Q18.2 Distributed execution

5 Points

Given your query, indicate where Spark would have to *repartition* or shard (also called *shuffle* or *exchange*) the data to perform the computation. Specify what the repartition key would be.

```
Shard on Enrolled.studentid, and on both group by columns:
Students.first and Enrolled.courseid.
```

## Q19 Costs

11 Points

Suppose we are given JSON document on people that looks like the following fragment.

```
[
  {
    "id": 1,
    "name": "Jun",
    "parent_ids": [],
    "child_ids": [4, 5]
  },
  {
    "id": 2,
    "name": "Maya",
    "parent_ids": [],
    "child_ids": [4, 5]
  },
  ...
  {
    "id": 4,
    "name": "Ava",
    "parent_ids": [1, 2],
    "child_ids": []
  },
]
```

### Q19.1 Schema for hierarchical data

5 Points

Propose a **schema** (with attribute types) for the above data. Use the syntax:

**Table(attrb1: type1, attrib2: type2, ...)**

and specify what would be the *keys* and *foreign keys* or references.

```
Person(id: int, name:string)
HasParent(parent_id: int, child_id: int)
```

Keys: id of person table

Foreign keys: parent\_id of HasParent, child\_id of HasParent

## Q19.2 Populating tables from JSON

6 Points

Using Pandas, Spark SQL, or the equivalent -- show how we would read the JSON data and populate the table(s).

```
schema = StructType([
  StructField("id", IntegerType()),
  StructField("name", StringType()),
  StructField("parent_ids", ArrayType()),
  StructField("child_ids", ArrayType())])

raw_data = spark.read.json("data.json", schema=schema)
raw_data.createOrReplaceTempView("data")
person = spark.sql("SELECT id, name FROM data")
hasParent = spark.sql("SELECT id as parent_id, col as child_id
FROM
(SELECT id, POSEXPLODE(child_ids) FROM data)
WHERE col IS NOT NULL ")
```

Midterm 1

● GRADED

STUDENT

Claudia Jiyun Zhu

TOTAL POINTS

**85.5 / 87 pts**

QUESTION 1

Directions

**6 / 6 pts**

QUESTION 2

Capturing data

**3 / 3 pts**

QUESTION 3

Query operations

**4 / 4 pts**

QUESTION 4

Knowledge representation

**3 / 3 pts**

QUESTION 5

NoSQL

**3 / 3 pts**

QUESTION 6

Integration

**3 / 3 pts**

QUESTION 7

Computer architecture

**3 / 3 pts**

QUESTION 8

Sharding

**4 / 4 pts**

QUESTION 9

Queries

**3 / 3 pts**

QUESTION 10

Statistical tests

**3 / 3 pts**

QUESTION 11

Indexing

**4 / 4 pts**

QUESTION 12

Breadth-first search

**3 / 3 pts**

**QUESTION 13**

Partitioning data

**3 / 3 pts****QUESTION 14**

Link analysis

**3 / 3 pts****QUESTION 15**

Matrices for machine learning

**3 / 3 pts****QUESTION 16**

Distributed processing

**6 / 6 pts**

16.1 Relational operator 1

**3 / 3 pts**

16.2 Relational operator 2

**3 / 3 pts****QUESTION 17**

Speedups

**8 / 8 pts**

17.1 Data structure 1

**4 / 4 pts**

17.2 Data structure 2

**4 / 4 pts****QUESTION 18**

Long-form query

**9.5 / 11 pts**

18.1 The query

**4.5 / 6 pts**

18.2 Distributed execution

**5 / 5 pts****QUESTION 19**

Costs

**11 / 11 pts**

19.1 Schema for hierarchical data

**5 / 5 pts**

19.2 Populating tables from JSON

**6 / 6 pts**